

Homework 4 - Neural Network vs LogReg

Soheil Sadeghi Eshkevari

07-06-2018

1 Introduction

In this report, results of the last homework is going to be presented. In this project, we trained our machine to classify MNIST dataset, which is a well-known dataset of digit handwritings. For this purpose, both Logistic Regression and Neural Network were tried. For each, three different solvers, Stochastic Gradient Decent (SGD), SGD with Momentum and ADAM were exploited. Since all these optimization techniques are stochastic, various batch sizes were tried and results compared. In the next sections, results will be presented.

2 Logistic Regression Results

The first classifier that was tried was the Logistic Regression. This method is basically a binary classifier, thus to be able to train our dataset using it, we turned our labels to binary. For this purpose, label 5 is assigned to be +1 and other labels to be -1. Results are shown below:

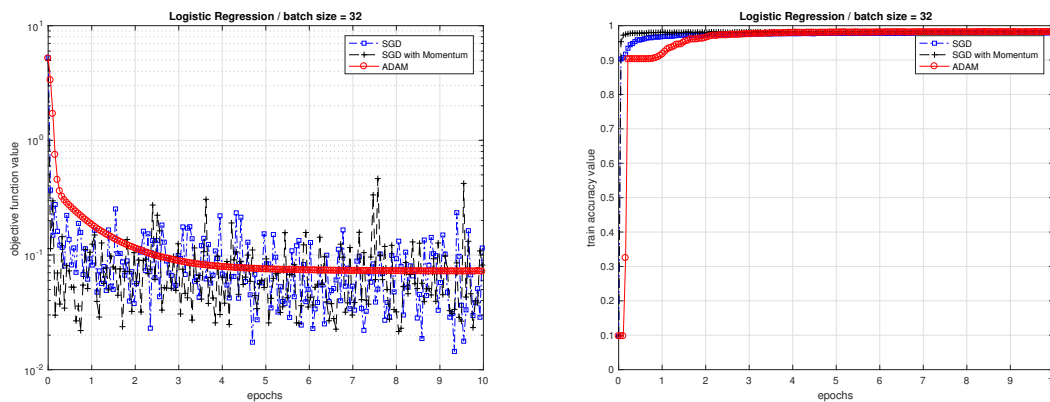


Figure 1: Performance plots vs. number of epochs with batch size 32

By comparing plots, following observations can be made. Less computational cost the problem needs when smaller batch sizes are set. In addition, with the setting I used for the training, ADAM is the most stable, while the least desirable one in terms of the fastest accuracy. Both SGD methods are very fast and suitable, however there show considerable fluctuations in their performance. This implies that the average of parameters can be a better alternative to get more stable results.

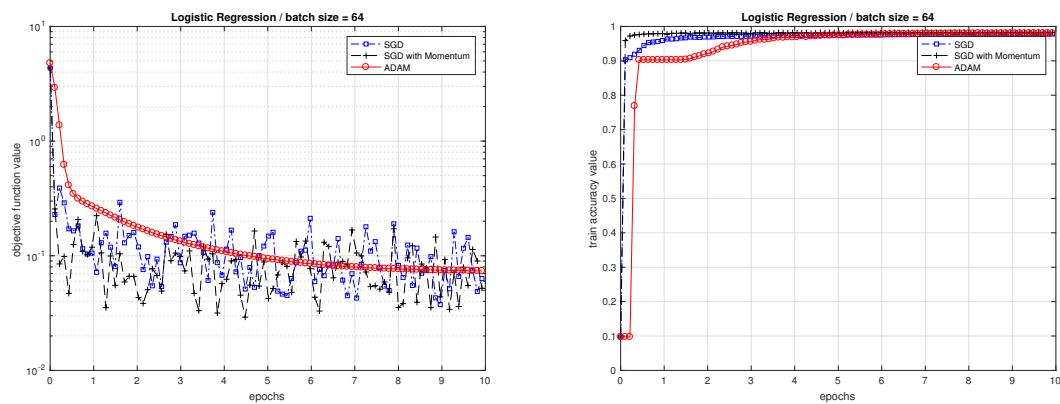


Figure 2: Performance plots vs. number of epochs with batch size 64




Figure 3: Performance plots vs. number of epochs with batch size 128

3 Neural Network Results, Multi-Label Classification

The next set of results belongs to the multi-class classification of the MNIST data using a 2 layer Neural Network. Sigmoid activation functions were used to connect layers. At the output layer, cross entropy function turns values to probabilities for each label. The number of nodes in the hidden layer also is set to be 40, based on both my own tuning and the suggested number from Mohammad and Mertcan's course project. As before, three different batch sizes were analyzed and results are shown below:

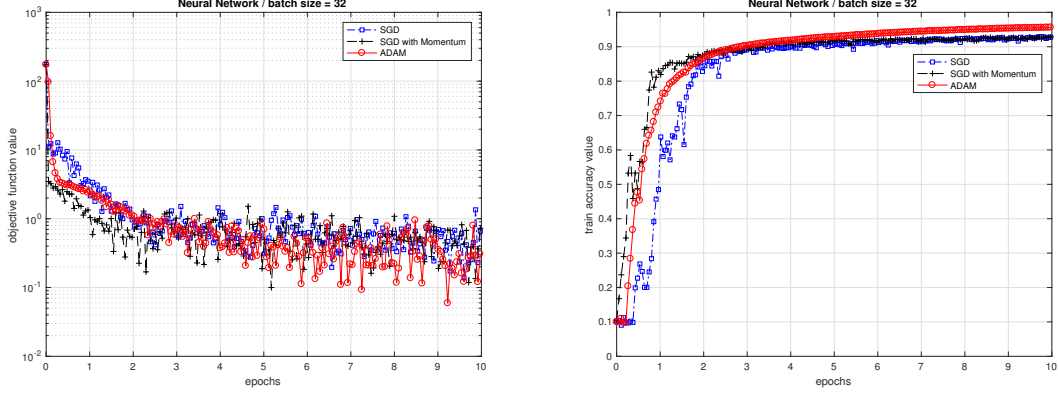


Figure 4: Performance plots vs. number of epochs with batch size 32

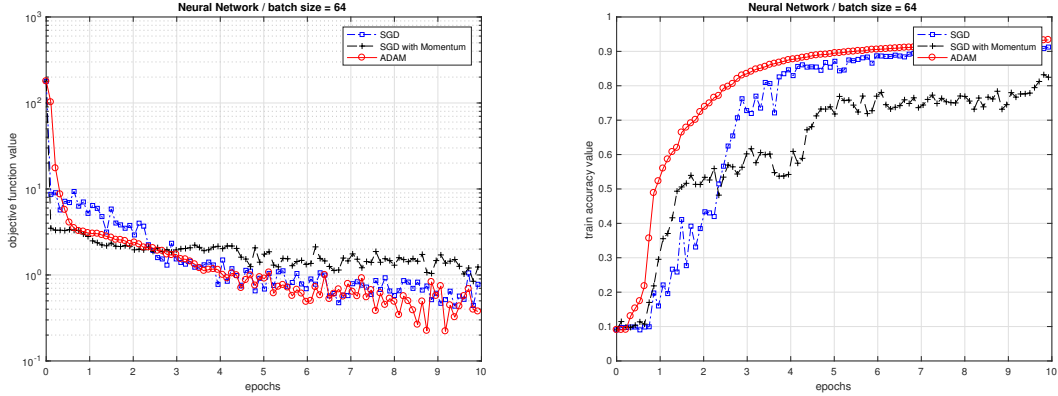


Figure 5: Performance plots vs. number of epochs with batch size 64

The most challenging part for NN training by SGD and SGD with momentum was to tune step sizes in a way that they converge to accurate results. In order to find the best values, first a fixed step was given and the best value for that was found (around 0.8). Next, a decaying step size in the format of $\frac{a}{b+k}$ was tuned to vary within 0.9 to 0.7. The step size was successful in the majority of cases, however, as it is noticeable in Figure 6, SGD with momentum in case of batch size equal to 128 was not able to get converged and resulted inaccurate predictions. In addition, by comparing different batch sizes, it can be interpreted that ADAM is obviously the dominant solver for NN since it is both stable and accurate. It also has no sensitivity to the step size and default settings worked appropriately.

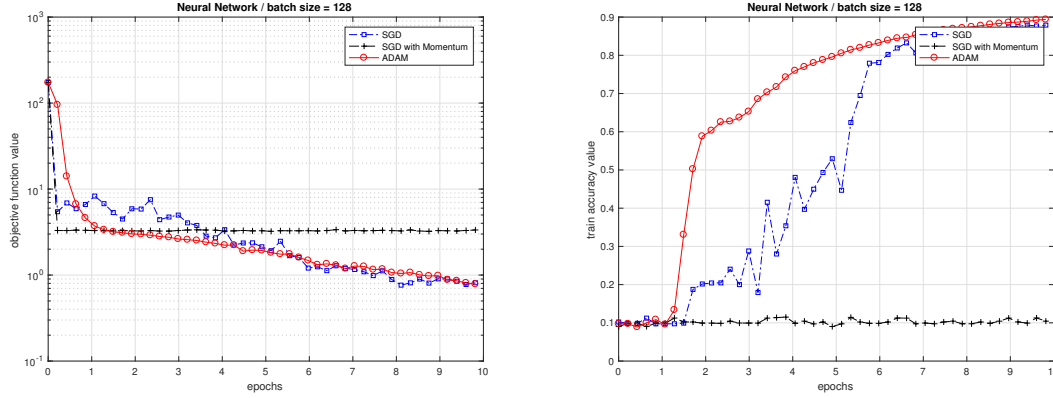


Figure 6: Performance plots vs. number of epochs with batch size 128

4 Fair comparison - Binary NN vs. Logistic Regression

As requested in the homework definition by Professor Scheinberg, in order to have a fair comparison between two methods, a binary classification problem with same logic explained for Logistic Regression was held using Neural Networks on MNIST dataset. The same analyses were done and results are presented hereafter:

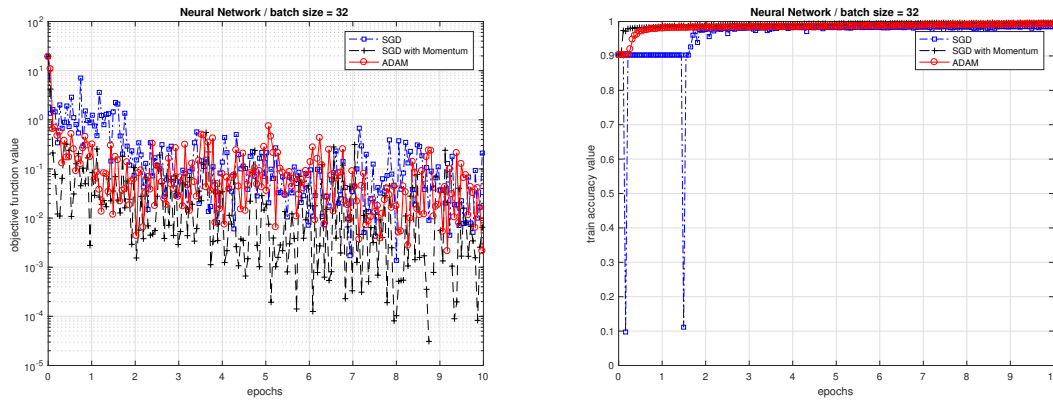


Figure 7: Performance plots vs. number of epochs with batch size 32

By comparing results, it can be deduced that NN is also very successful for binary classification of MNIST. The best accuracy of NN is slightly higher than the best accuracy derived from Logistic Regression. However, as before, for higher batch sizes, less desired behavior can be observed by NN. The most desired setting between these two methods is to use NN with batch size equal to 32 and using ADAM or SGD with momentum as the solver.

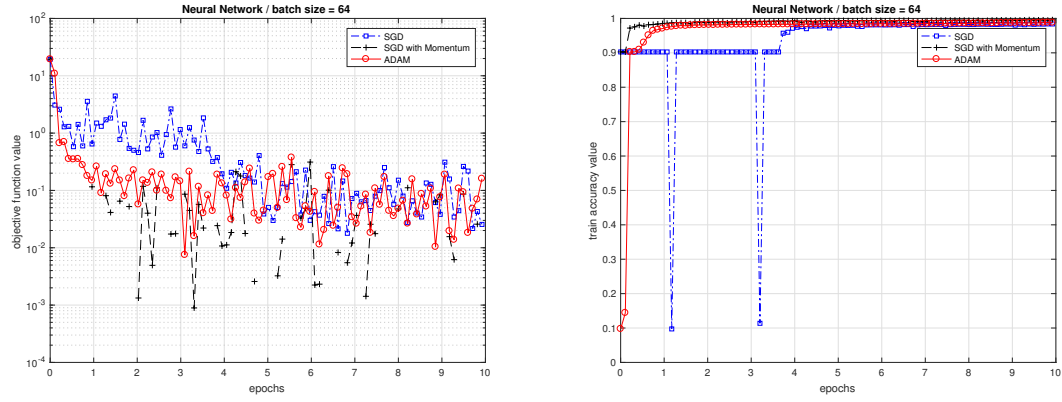


Figure 8: Performance plots vs. number of epochs with batch size 64




Figure 9: Performance plots vs. number of epochs with batch size 128

5 Code

In this section, the main body of my NN code for this homework is presented. However, subroutines and associated files will be sent in a package. Please note that a similar file exists for Logistic Regression and can be found in the package.

```
1 % DATABASE 4: mnist.mat (always run this section in advance, for all for four algorithms)
2
3 clc;
4 clear;
5
6 rng shuffle
7
8 data_1 = load('new_data/mnist_mult.mat');
9
10 X_temp = data_1.TrainX;
11 Y = data_1.TrainY;
12 Xt_temp = data_1.TestX;
13 Yt = data_1.TestY;
14
15 % fair comparison (Binary class NN) -----
16 Y_temp = zeros(size(Y,1),2);
17 Y_temp(:,1) = Y(:,5)==1;
18 Y_temp(:,2) = Y(:,5)==0;
19 Y = Y_temp;
20
21 Y_temp = zeros(size(Yt,1),2);
22 Y_temp(:,1) = Yt(:,5)==1;
23 Y_temp(:,2) = Yt(:,5)==0;
24 Yt = Y_temp;
25 Y_temp = [];
26 % -----
27
28 X = zeros(60000,28*28);
29 for i = 1:60000
30     temp = squeeze(X_temp(:,:,i));
31     X(i,:) = temp(:);
32 end
33
34 Xt = zeros(10000,28*28);
35 for i = 1:10000
36     temp = squeeze(Xt_temp(:,:,i));
37     Xt(i,:) = temp(:);
38 end
39
40 %% Scaling and parameters
41
42 for i = 1:size(X,1)
43     X(i,:) = X(i,:)/norm(X(i,:),2);
44 end
45 for i = 1:size(Xt,1)
46     Xt(i,:) = Xt(i,:)/norm(Xt(i,:),2);
47 end
48
49 input_layer_size = 28*28;% 28x28 Input Images of Digits
50 hidden_layer_size = 40;% 40 hidden units
51 num_labels = 2;          % 10 labels, from 1 to 10
52                        % (note that we have mapped "0" to label 10)
53 lambda = 0.0;
54
55 % w_1 = rand(hidden_layer_size, (input_layer_size + 1));
56 % w_2 = rand(num_labels, (hidden_layer_size + 1));
```

```

57 w_1 = rand(hidden_layer_size, (input_layer_size));
58 w_2 = rand(num_labels, (hidden_layer_size));
59 nn_params = [w_1(:); w_2(:)];
60 % a_2 = 1./(1+X'*w_1);
61 % p = 1./(1+a_2*w_2);
62 % hs = -log(exp(p)./sum(exp(p),2));
63
64 batch_size = 128; % default = 50
65 epochs = 10;
66 iter = ceil(epochs*size(X,1)/batch_size);
67
68 %% solver (ADAM)
69
70 alfa = 0.002; % 0.002
71 beta_1 = 0.9;
72 beta_2 = 0.999;
73 eps = 1e-8;
74 m = zeros(numel(nn_params),1);
75 v = zeros(numel(nn_params),1);
76 t = 0;
77 thresh = 1e-2;
78
79 grad = 1;
80 J_tr_ADAM = zeros(iter,1);
81 J_ts_ADAM = zeros(iter,1);
82 nn_params_ADAM = zeros(iter+1,numel(nn_params));
83 nn_params_ADAM(1,:) = nn_params;
84
85 for j = 1:iter
86     t = t + 1;
87     Sk = randi(size(X,1),1,batch_size);
88     [J_tr_ADAM(j), grad] = nnCostFunction(nn_params_ADAM(j,:), ...
89                                         input_layer_size, ...
90                                         hidden_layer_size, ...
91                                         num_labels, ...
92                                         X(Sk,:), Y(Sk,:), lambda);
93     m = beta_1*m + (1-beta_1)*grad;
94     v = beta_2*v + (1-beta_2)*(grad.^2);
95     m_hat = m/(1-beta_1^t);
96     v_hat = v/(1-beta_2^t);
97     nn_params_ADAM(j+1,:) = nn_params_ADAM(j,:) - (alfa * m_hat./(sqrt(v_hat)+eps))';
98     J_ts_ADAM(j) = nnCost(nn_params_ADAM(j+1,:), ...
99                           input_layer_size, ...
100                          hidden_layer_size, ...
101                          num_labels, ...
102                          X_t, ...
103                          Y_t, lambda);
104
105     % obj_fnc = J_tr_ADAM(j)
106     % norm(grad,2)
107 end
108
109 % plot(1:20:iter,[J_tr_ADAM(1:20:end) J_ts_ADAM(1:20:end)]);
110
111 %% solver (SGD)
112
113 a = 28350;
114 b = 30500;
115 teta = 0.95;
116
117 k = 1:iter;
118 nu = a./(b+k);
119
120 J_tr_SGD = zeros(iter,1);
121 J_ts_SGD = zeros(iter,1);

```

```

122 nn_params_SGD = zeros(iter+1,numel(nn_params));
123 nn_params_SGD(1,:) = nn_params;
124
125 v = zeros(length(nn_params),1);
126 v_prev = zeros(length(nn_params),1);
127
128 for j = 1:iter
129     Sk = randi(size(X,1),1,batch_size);
130     [J_tr_SGD(j), grad] = nnCostFunction(nn_params_SGD(j,:), ...
131                                         input_layer_size, ...
132                                         hidden_layer_size, ...
133                                         num_labels, ...
134                                         X(Sk,:), Y(Sk,:), lambda);
135     v = teta * v_prev + grad;
136     v_prev = v;
137     nn_params_SGD(j+1,:) = nn_params_SGD(j,:) - nu(j) * grad';
138     % nn_params = nn_params - nu(j) * grad;
139     % nn_params = nn_params - 0.1 * v;
140     J_ts_SGD(j) = nnCost(nn_params_SGD(j+1,:), ...
141                         input_layer_size, ...
142                         hidden_layer_size, ...
143                         num_labels, ...
144                         X_t, ...
145                         Y_t, lambda);
146     j
147     % J_tr_SGD(j)
148     % norm(grad,2)
149 end
150
151 % plot(J_ts_SGD(2:j)); hold on; plot(J_tr_SGD(2:j));
152
153 %% solver (SGD + Momentum)
154
155 a = 28350;
156 b = 30500;
157 teta = 0.95;
158
159 k = 1:iter;
160 nu = a./(b+k);
161
162 J_tr_SGDM = zeros(iter,1);
163 J_ts_SGDM = zeros(iter,1);
164 nn_params_SGDM = zeros(iter+1,numel(nn_params));
165 nn_params_SGDM(1,:) = nn_params;
166 v = zeros(length(nn_params),1);
167 v_prev = zeros(length(nn_params),1);
168
169 for j = 1:iter
170     Sk = randi(size(X,1),1,batch_size);
171     [J_tr_SGDM(j), grad] = nnCostFunction(nn_params_SGDM(j,:), ...
172                                         input_layer_size, ...
173                                         hidden_layer_size, ...
174                                         num_labels, ...
175                                         X(Sk,:), Y(Sk,:), lambda);
176     v = teta * v_prev + grad;
177     v_prev = v;
178     % nn_params = nn_params - nu(j) * grad;
179     % nn_params = nn_params - nu(j) * grad;
180     nn_params_SGDM(j+1,:) = nn_params_SGDM(j,:) - nu(j) * v';
181     J_ts_SGDM(j) = nnCost(nn_params_SGDM(j+1,:), ...
182                         input_layer_size, ...
183                         hidden_layer_size, ...
184                         num_labels, ...
185                         X_t, ...
186                         Y_t, lambda);

```



```

187                                     j
188 %     J_tr_SGDM(j)
189 %     norm(grad,2)
190 end
191
192 % plot(J_ts_SGDM(2:j)); hold on; plot(J_tr_SGDM(2:j));
193
194 %% plots
195 step = 100;
196 % function in three methods
197 semilogy((1:step:iter)*batch_size/size(X,1),J_tr_SGD(1:step:end), '-.bs'); hold on
198 semilogy((1:step:iter)*batch_size/size(X,1),J_tr_SGDM(1:step:end), '—k+'); hold on
199 semilogy((1:step:iter)*batch_size/size(X,1),J_tr_ADAM(1:step:end), '-ro');
200 grid on
201 legend('SGD','SGD with Momentum','ADAM');
202 xlabel('epochs');
203 ylabel('objective function value');
204 title(['Neural Network / batch size = ',num2str(batch_size)])
205
206 % accuracy in three methods
207 y_vec = zeros(size(Y,1),1);
208 for i = 1:size(Y,1)
209     y_vec(i) = find(Y(i,:)==1);
210 end
211
212 pred_SGD = zeros(1,length(1:step:iter));
213 pred_SGDM = zeros(1,length(1:step:iter));
214 pred_ADAM = zeros(1,length(1:step:iter));
215 for i = 1:step:iter
216     pred_SGD(i) = mean(double(predict(nn_params_SGD(i+1,:), ...
217                                     X,hidden_layer_size,input_layer_size,num_labels) == y_vec));
218     pred_SGDM(i) = mean(double(predict(nn_params_SGDM(i+1,:), ...
219                                     X,hidden_layer_size,input_layer_size,num_labels) == y_vec));
220     pred_ADAM(i) = mean(double(predict(nn_params_ADAM(i+1,:), ...
221                                     X,hidden_layer_size,input_layer_size,num_labels) == y_vec));
222 end
223 figure;
224 plot((1:step:iter)*batch_size/size(X,1),pred_SGD(1:step:end), '-.bs'); hold on
225 plot((1:step:iter)*batch_size/size(X,1),pred_SGDM(1:step:end), '—k+'); hold on
226 plot((1:step:iter)*batch_size/size(X,1),pred_ADAM(1:step:end), '-ro');
227 grid on
228 legend('SGD','SGD with Momentum','ADAM');
229 xlabel('epochs');
230 ylabel('train accuracy value');
231 title(['Neural Network / batch size = ',num2str(batch_size)])

```

6 Conclusions

In this homework, the main catch was to have a hands-on experience on coding some *off the shelf* methods and solvers and try to play with parameters to get a deeper understanding about these, rather than just superficially use them as blackbox toolsuites.